



Obesity Data Analysis Project - Predicting Obesity

Peter Gray

September 4, 2025

Produced with Quarto and Python

Table of contents

List of Python and R Packages	3
Data Loading and Wrangling	4
Data Wrangling and BMI Age Group Variable Creation	4
Logistic Regression	4
Variable Selection	4
Remove the NAs	4
Creation of Training Data	5
Logistic Regression Model 1	5
Coefficients Logistic Regression Model 1	5
Confusion Matrix - Logistic Regression Model 1	7
Classification Report Logistic Regression Model	8
LASSO Rgression	10
Discarded Features	12
Comparison of Logistc and LASSO ROC curve	13
Random Forest Plot	14
Comparison of all 3 models	16
Comparison of Logistc, Radnom Forst and LASSO ROC curve	16
Models removing colinearity	17
Colinearity	17
Identifying Columns to Drop	17
Remove the NAs	18
Confusion Matrix - Logistic Regression Model 2	18
Classification Report Logistic Regression Model 2	20
LASSO Rgression	22
Discarded Features	24
Comparison of Logistc and LASSO ROC curve	25
Random Forest Plot	26
Comparison of all 3 models	28
Comparison of Logistc and LASSO ROC curve	28
Conclusion from Logistic Model 1	30
Run the model agin to make sure we are not confused	30
Odds Ratios	30

List of Python and R Packages

Python

NumPy
Pandas
Matplotlib
seaborn
sklearn
statsmodels.api

Data Loading and Wrangling

Master dataframe is created in the data analysis project href [here](#)

```
import pandas as pd
import numpy as np

df =
    ↪ pd.read_csv(r"/home/pgr16/Documents/Data_Analysis/Obesity_Analysis_and_Prediction/Data/Master
```

Data Wrangling and BMI Age Group Variable Creation

```
df['HEIGHT'] = df['HEIGHT']/100 # convert to meteres

df['BMI'] = df['WEIGHT'] / (df['HEIGHT']**2)
df['Obese_Y_N'] = np.where(df['BMI'] >= 30, 1, 0)
```

Logistic Reegression

Variable Selection

```
y_var = df["Obese_Y_N"] # outcome
X_var = df[["AGE_GROUP", "GENDER", "RACE", "BIRTH_COUNTRY",
            "SALT_IN_PREP", "NUMBER_FOOD_DRINKS", "CALORIES",
            "NUT_CARBS_GRAMS", "NUT_FAT_GRAMS",
            "NUT_CAFFEINE_MILIGRAMS", "NUT_ALCOHOL_GRAMS",
            "TOTAL_WATER_DRANK_LITRE"]]

X_var = pd.get_dummies(X_var, drop_first=True)

# Drop the NAs
```

Remove the NAs

```
X_var = X_var.dropna()
y_var = y_var.loc[X_var.index]
```

Creation of Training Data

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_var, y_var,
    ↪ test_size=0.3, random_state=42)
```

Logistic Regression Model 1

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

logmod = LogisticRegression(max_iter=10000, random_state=0)

logmod.fit(X_train, y_train)

acc = accuracy_score(y_test, logmod.predict(X_test)) *100

print(f"Logistic Regression model accuracy: {acc:.2f}%")
```

Logistic Regression model accuracy: 68.09%

Coefficients Logistic Regression Model 1

```
coefficients = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': logmod.coef_[0]
})

# Optional: Sort by absolute value of coefficients
coefficients['Abs_Coefficient'] = coefficients['Coefficient'].abs()
coefficients = coefficients.sort_values(by='Abs_Coefficient', ascending=False)
coefficients
```

	Feature	Coefficient	Abs_Coefficient
9	AGE_GROUP_Less than 18	-2.000270	2.000270
14	RACE_Other Race - Including Multi-Racial	-0.908169	0.908169
12	RACE_Non-Hispanic White	-0.585536	0.585536
18	SALT_IN_PREP_Rarely	0.571331	0.571331
17	SALT_IN_PREP_Occasionally	0.558330	0.558330

	Feature	Coefficient	Abs_Coefficient
7	AGE_GROUP_36-60	0.514841	0.514841
16	SALT_IN_PREP_Never	0.445018	0.445018
19	SALT_IN_PREP_Very often	0.433826	0.433826
15	BIRTH_COUNTRY_US	0.388024	0.388024
8	AGE_GROUP_Above 60	0.382325	0.382325
10	GENDER_Male	-0.213255	0.213255
13	RACE_Other Hispanic	-0.209125	0.209125
11	RACE_Non-Hispanic Black	-0.122874	0.122874
6	TOTAL_WATER_DRANK_LITRE	0.066987	0.066987
0	NUMBER_FOOD_DRINKS	-0.054301	0.054301
5	NUT_ALCOHOL_GRAMS	-0.003438	0.003438
3	NUT_FAT_GRAMS	0.002851	0.002851
2	NUT_CARBS_GRAMS	-0.001171	0.001171
4	NUT_CAFFEINE_MILIGRAMS	0.000537	0.000537
1	CALORIES	0.000073	0.000073

Confusion Matrix - Logistic Regression Model 1

```
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Predictions
y_pred_train = logmod.predict(X_train)

conf_matrix_train = confusion_matrix(y_train, y_pred_train)
conf_matrix_train

sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




Figure 1: Confusion Matrix for Logistic Regression Model 1

Classification Report Logistic Regression Model

```
print(classification_report(y_train, y_pred_train))
```

	precision	recall	f1-score	support
0	0.72	0.89	0.79	3174
1	0.53	0.27	0.35	1511
accuracy			0.69	4685
macro avg	0.63	0.58	0.57	4685
weighted avg	0.66	0.69	0.65	4685

```
from sklearn.metrics import roc_curve, roc_auc_score, RocCurveDisplay

y_proba_test = logmod.predict_proba(X_test)[:, 1]

fpr, tpr, _ = roc_curve(y_test, y_proba_test)
auc_score = roc_auc_score(y_test, y_proba_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")

plt.plot([0, 1], [0, 1], "k--", label="Random")

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```





Figure 2: ROC Curve for Logistic Regression Model 1

Since the Model is OK, but could we better? We will try a LASSO regression and a Random Forest Plot to see if we can get a better model.

LASSO Rgression

```
lasso = LogisticRegression(penalty = 'l1', solver = 'liblinear',
    ↪ class_weight='balanced', C=1.0)
lasso.fit(X_train, y_train)

coefs = lasso.coef_[0] # Get coefficients
feature_names = X_train.columns

coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefs})

kept = coef_df[coef_df['Coefficient'] != 0].sort_values(by='Coefficient',
    ↪ key=abs, ascending=False)
discarded = coef_df[coef_df['Coefficient'] == 0]

report_dict = classification_report(y_test, lasso.predict(X_test),
    ↪ output_dict=True)
class_report = pd.DataFrame(report_dict).transpose()
```

```
y_proba_lasso_test = lasso.predict_proba(X_test)[: , 1]
fpr_lasso, tpr_lasso, _ = roc_curve(y_test, y_proba_lasso_test)
auc_lasso = roc_auc_score(y_test, y_proba_lasso_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_lasso, tpr_lasso, label=f"Logistic Regression (AUC =
    ↪ {auc_lasso:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random')

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 3: ROC Curve for the Logistic Regression Model 1

Discarded Features

```
plt.figure(figsize=(10, 6))
plt.barh(kept['Feature'], kept['Coefficient'], color='darkorange')
plt.xlabel("Coefficient Value")
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```




Figure 4: LASSO-Selected Feature Coefficients

Comparison of Logistic and LASSO ROC curve

```
plt.figure(figsize=(8, 6))
plt.plot(fpr_lasso, tpr_lasso, label=f"LASSO Logistic (AUC = {auc_lasso:.2f})")
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```

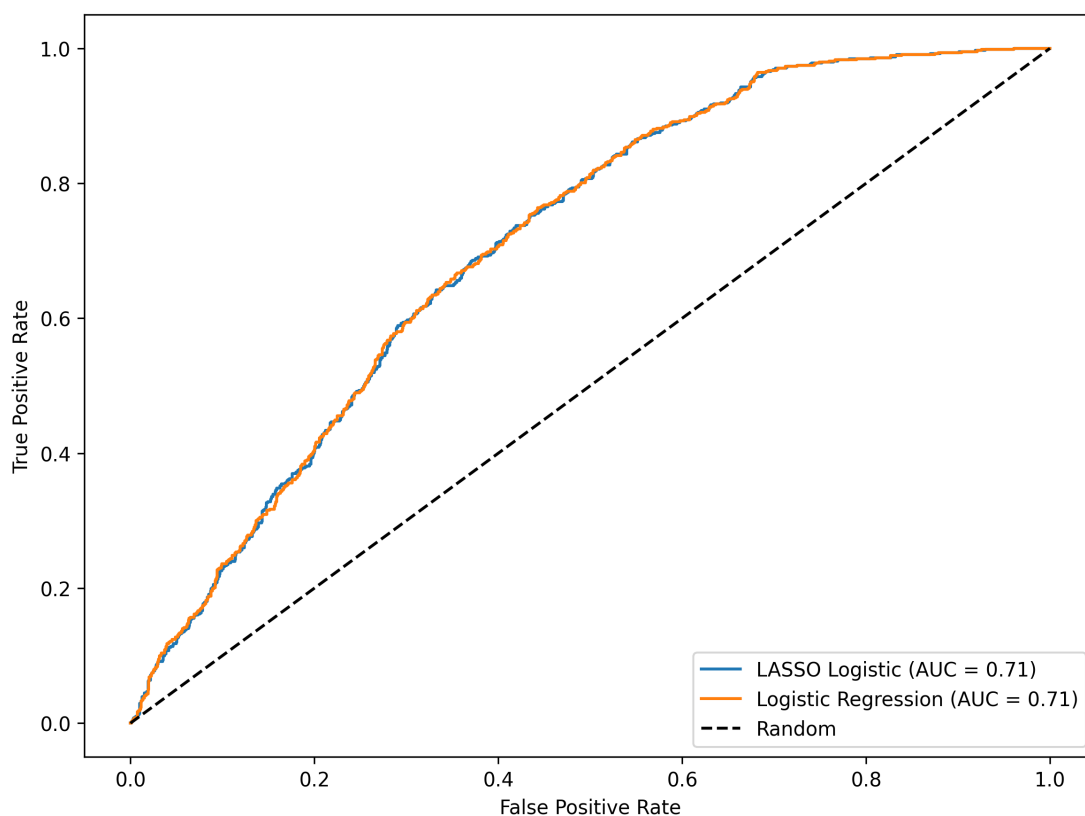


Figure 5: ROC Curve comparison between LASSO and Logistic Regression Model

Random Forest Plot

```
from sklearn.ensemble import RandomForestClassifier as randforest

rf = randforest(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

y_proba_randforest_test = rf.predict_proba(X_test)[: , 1]
fpr_randforest, tpr_randforest, _ = roc_curve(y_test, y_proba_randforest_test)
auc_randforest = roc_auc_score(y_test, y_proba_randforest_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_randforest, tpr_randforest, label=f"Random Forest (AUC =
↪ {auc_randforest:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 6: ROC Curve for Random Forest Model

Comparison of all 3 models

Comparison of Logistic, Random Forest and LASSO ROC curve

```
plt.figure(figsize=(8, 6))
plt.plot(
    fpr_randforest, tpr_randforest, label=f"Random Forest (AUC = {auc_randforest:.2f})"
)
plt.plot(fpr_lasso, tpr_lasso, label=f"LASSO Logistic (AUC = {auc_lasso:.2f})")
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")
plt.plot([0, 1], [0, 1], "k--", label="Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```

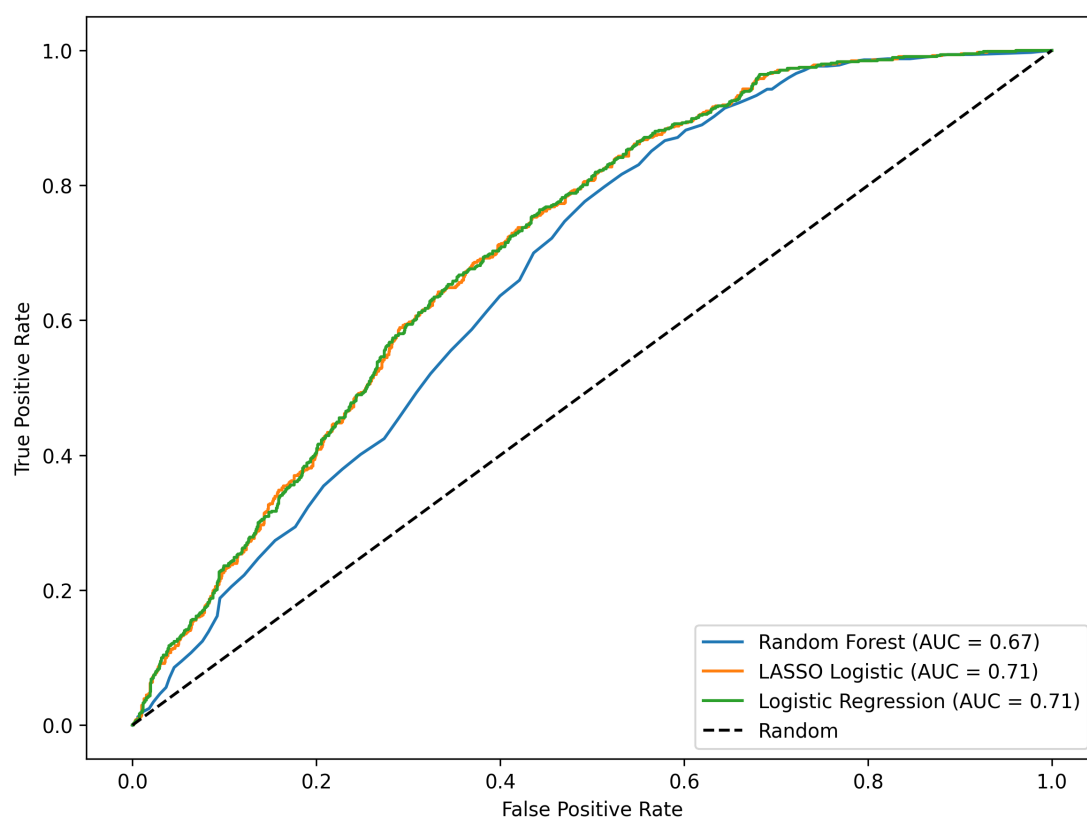


Figure 7: ROC Curve comparison between LASSO and Logistic Regression Model and Random Forest Model

Models removing colinearity

Colinearity

```
corr = X_var.corr()  
plt.figure(figsize=(12,8))  
sns.heatmap(corr, annot=False, cmap="coolwarm", center=0)  
plt.show()
```




Figure 8: Correlation Matrix of Predictors

Identifying Columns to Drop

```
upper_tri = corr.where(np.triu(np.ones(corr.shape), k=1).astype(np.bool))
```

```
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] >
↪ 0.85)]
```

```
print(to_drop)
```

```
['NUT_CARBS_GRAMS', 'NUT_FAT_GRAMS']
```

'NUT_CARBS_GRAMS' and 'NUT_FAT_GRAMS' are identified as highly colinear (>0.85) and therefore we will remove them to see if there is any improvement in the models we run before

```
df2 = df.drop(to_drop, axis = 1)
```

```
y_var = df2["Obese_Y_N"] # outcome
X_var = df2[["AGE_GROUP", "GENDER", "RACE", "BIRTH_COUNTRY",
            "SALT_IN_PREP", "NUMBER_FOOD_DRINKS", "CALORIES",
            "NUT_CAFFEINE_MILIGRAMS", "NUT_ALCOHOL_GRAMS",
            "TOTAL_WATER_DRANK_LITRE"]]
```

```
X_var = pd.get_dummies(X_var, drop_first=True)
```

```
# Drop the NAs
```

Remove the NAs

```
X_var = X_var.dropna()
y_var = y_var.loc[X_var.index]
```

Confusion Matrix - Logistic Regression Model 2

```
from sklearn.metrics import classification_report, confusion_matrix

# Predictions
y_pred_train = logmod.predict(X_train)

conf_matrix_train = confusion_matrix(y_train, y_pred_train)
conf_matrix_train

sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




Figure 9: Confusion Matrix Logistic Regression Model 2

Classification Report Logistic Regression Model 2

```
print(classification_report(y_train, y_pred_train))
```

	precision	recall	f1-score	support
0	0.72	0.89	0.79	3174
1	0.53	0.27	0.35	1511
accuracy			0.69	4685
macro avg	0.63	0.58	0.57	4685
weighted avg	0.66	0.69	0.65	4685

```
from sklearn.metrics import roc_curve, roc_auc_score, RocCurveDisplay

y_proba_test = logmod.predict_proba(X_test)[: , 1]

fpr, tpr, _ = roc_curve(y_test, y_proba_test)
auc_score = roc_auc_score(y_test, y_proba_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random')

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```

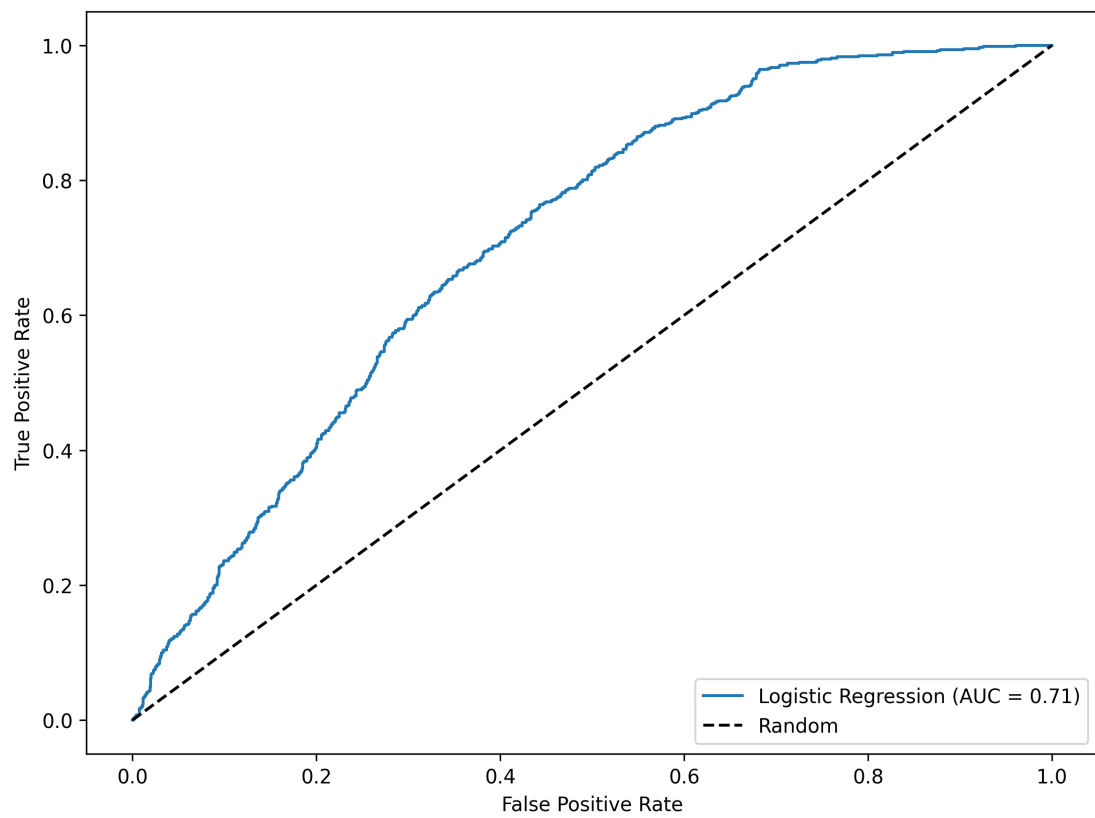


Figure 10: ROC Curve Logistic Regression Model 2

Since the Model is OK, but could we better? We will try a LASSO regression and a Random Forest Plot to see if we can get a better model.

LASSO Rgression

```
lasso = LogisticRegression(penalty = 'l1', solver = 'liblinear',
    ↪ class_weight='balanced', C=1.0)
lasso.fit(X_train, y_train)

coefs = lasso.coef_[0] # Get coefficients
feature_names = X_train.columns

coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefs})

kept = coef_df[coef_df['Coefficient'] != 0].sort_values(by='Coefficient',
    ↪ key=abs, ascending=False)
discarded = coef_df[coef_df['Coefficient'] == 0]

report_dict = classification_report(y_test, lasso.predict(X_test),
    ↪ output_dict=True)
class_report = pd.DataFrame(report_dict).transpose()
```

```
y_proba_lasso_test = lasso.predict_proba(X_test)[: , 1]
fpr_lasso, tpr_lasso, _ = roc_curve(y_test, y_proba_lasso_test)
auc_lasso = roc_auc_score(y_test, y_proba_lasso_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_lasso, tpr_lasso, label=f"Logistic Regression (AUC =
    ↪ {auc_lasso:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random')

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 11: ROC Curve LASSO Regression Model 2

Discarded Features

```
plt.figure(figsize=(10, 6))
plt.barh(kept['Feature'], kept['Coefficient'], color='darkorange')
plt.xlabel("Coefficient Value")
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```

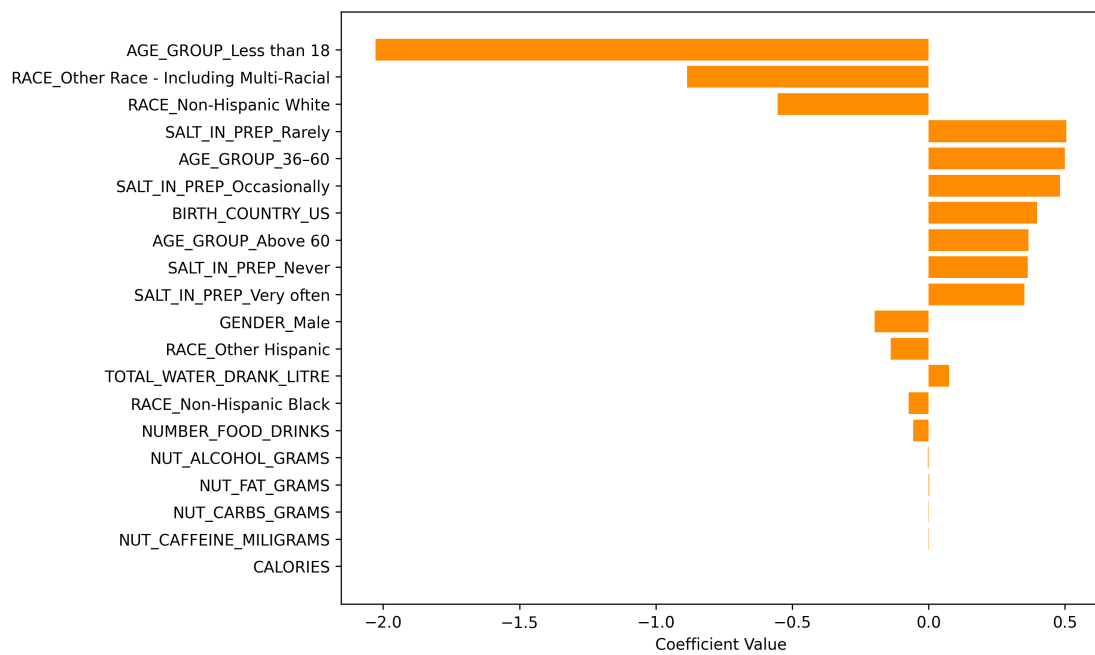


Figure 12: Discarded Features Regression Model 2

Comparison of Logistic and LASSO ROC curve

```
plt.figure(figsize=(8, 6))
plt.plot(fpr_lasso, tpr_lasso, label=f"LASSO Logistic (AUC = {auc_lasso:.2f})")
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 13: ROC Curve LASSO Logistic Regression

Random Forest Plot

```
from sklearn.ensemble import RandomForestClassifier as randforest

rf = randforest(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

y_proba_randforest_test = rf.predict_proba(X_test)[:, 1]
fpr_randforest, tpr_randforest, _ = roc_curve(y_test, y_proba_randforest_test)
auc_randforest = roc_auc_score(y_test, y_proba_randforest_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_randforest, tpr_randforest, label=f"Random Forest (AUC =
↪ {auc_randforest:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 14: ROC Curve Random Forest Graph

Comparison of all 3 models

Comparison of Logistic and LASSO ROC curve

```
plt.figure(figsize=(8, 6))
plt.plot(fpr_randforest, tpr_randforest, label=f"Random Forest (AUC = {auc_randforest:.2f})")
plt.plot(fpr_lasso, tpr_lasso, label=f"LASSO Logistic (AUC = {auc_lasso:.2f})")
plt.plot(fpr, tpr, label=f"Logistic Regression (AUC = {auc_score:.2f})")
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(False)
plt.tight_layout()
plt.show()
```




Figure 15: ROC Curve, LASSO, Random Forest Logistic Regression

There is no no real difference in the models that have highly colinear removed. We will revert

back to the original Logistic Regression Model with all identified .

Conclusion from Logistic Model 1

Run the model again to make sure we are not confused

We will need to provide confidence intervals. We'll do this by running the statsmodel.api rather than sci-kit-learn

```
import statsmodels.api as sm
y_var = df["Obese_Y_N"] # outcome
X_var = df[["AGE_GROUP", "GENDER", "RACE", "BIRTH_COUNTRY",
            "SALT_IN_PREP", "NUMBER_FOOD_DRINKS", "CALORIES",
            "NUT_CARBS_GRAMS", "NUT_FAT_GRAMS",
            "NUT_CAFFEINE_MILIGRAMS", "NUT_ALCOHOL_GRAMS",
            "TOTAL_WATER_DRANK_LITRE"]]

X_var = pd.get_dummies(X_var, drop_first=True)
X_var = X_var.dropna()
y_var = y_var.loc[X_var.index]

X_sm = sm.add_constant(X_var).astype(float)
y_sm = y_var.astype(float)

# Fit logistic regression
logit_model = sm.Logit(y_var, X_sm)
result = logit_model.fit()
```

```
Optimization terminated successfully.
      Current function value: 0.548471
      Iterations 7
```

Odds Ratios

```
params = result.params
conf = result.conf_int()
conf.columns = ["2.5%", "97.5%"]

or_table = pd.DataFrame({
    "Feature": params.index,
    "Coefficient": params,
    "Odds Ratio": np.exp(params),
    "2.5%": np.exp(conf["2.5%"]),
    "97.5%": np.exp(conf["97.5%"]),
})
```

```

    "p-value": result.pvalues})

or_table['p-value'] = np.where(or_table['p-value'] < 0.05, "*p <0.05",
    ↳ or_table['p-value'])
or_table = or_table[or_table["Feature"] != "const"]

or_table["Feature"] = or_table["Feature"].str.replace("_", " ",
    ↳ regex=False).str.title()

or_table.sort_values(by="Odds Ratio", ascending = False)

```

	Feature	Coefficient	Odds Ratio
AGE_GROUP_36–60	Age Group 36–60	0.639804	1.896110
SALT_IN_PREP_Rarely	Salt In Prep Rarely	0.560211	1.751042
AGE_GROUP_Above 60	Age Group Above 60	0.510549	1.666206
SALT_IN_PREP_Occasionally	Salt In Prep Occasionally	0.495983	1.642112
SALT_IN_PREP_Very often	Salt In Prep Very Often	0.430922	1.538675
BIRTH_COUNTRY_US	Birth Country Us	0.412340	1.510348
SALT_IN_PREP_Never	Salt In Prep Never	0.365931	1.441856
TOTAL_WATER_DRANK_LITRE	Total Water Drank Litre	0.067259	1.069573
NUT_FAT_GRAMS	Nut Fat Grams	0.002877	1.002881
NUT_CAFFEINE_MILIGRAMS	Nut Caffeine Miligrams	0.000389	1.000389
CALORIES	Calories	0.000059	1.000059
NUT_CARBS_GRAMS	Nut Carbs Grams	-0.000989	0.999012
NUT_ALCOHOL_GRAMS	Nut Alcohol Grams	-0.002111	0.997891
NUMBER_FOOD_DRINKS	Number Food Drinks	-0.051876	0.949446
RACE_Non-Hispanic Black	Race Non-Hispanic Black	-0.159616	0.852471
GENDER_Male	Gender Male	-0.194653	0.823120
RACE_Other Hispanic	Race Other Hispanic	-0.288626	0.749292
RACE_Non-Hispanic White	Race Non-Hispanic White	-0.616849	0.539642
RACE_Other Race - Including Multi-Racial	Race Other Race - Including Multi-Racial	-0.844126	0.429933
AGE_GROUP_Less than 18	Age Group Less Than 18	-1.894983	0.150321

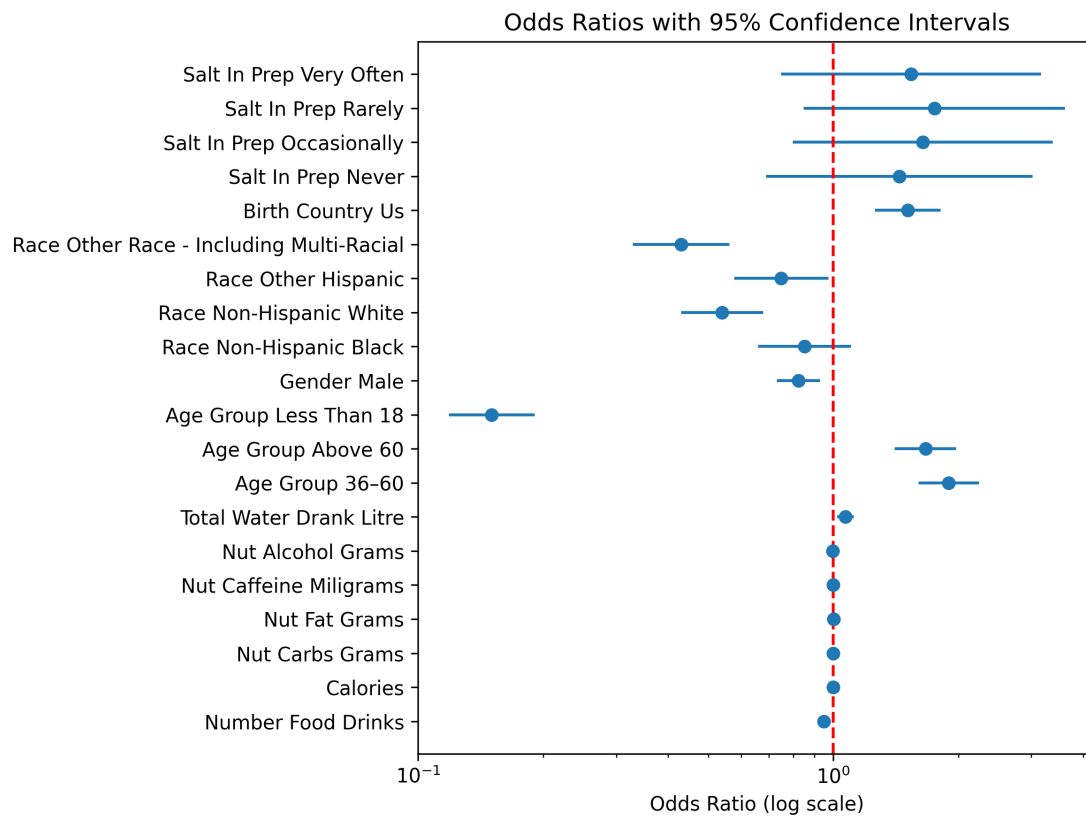
Accordng to our model Calories alone do not necessarily appear to predict whether a person is going to be obese or not.

Forest Plot

```
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6))
plt.errorbar(or_table["Odds Ratio"], or_table["Feature"],
             xerr=[or_table["Odds Ratio"] - or_table["2.5%"],
                   or_table["97.5%"] - or_table["Odds Ratio"]],
             fmt='o')
plt.axvline(x=1, color="red", linestyle="--")
plt.xscale("log")
plt.xlabel("Odds Ratio (log scale)")
plt.title("Odds Ratios with 95% Confidence Intervals")
plt.tight_layout()
plt.figtext(
    0.5, -0.05,
    "Odds Ratio > 1 → Higher likelihood of Obesity | Odds Ratio < 1 → Lower  

    ↳ likelihood of Obesity",
    ha="center", fontsize=10, style="italic"
)
plt.show()
```

Odds Ratio > 1 → Higher likelihood of Obesity | Odds Ratio < 1 → Lower likelihood of Obesity

Figure 16: FForest Plot of Odds Ratios for Logistic Regression Model 1